

Supervised classification for driver types with aggregation analysis.

```
In [ ]: import pandas as pd
```

```
In [2]: fcd_df = pd.read_excel("C:/Users/saiko/OneDrive/Desktop/699/Task-7/fcd-output
emission_df = pd.read_excel("C:/Users/saiko/OneDrive/Desktop/699/Task-7/emiss:
features = pd.read_excel("C:/Users/saiko/OneDrive/Desktop/699/Task-7/features
```

```
In [28]: features
```

```
Out[28]:
```

	Unnamed: 0	time	id	x	y	angle	type	speed	pos	lane
0	0	0.0	f_0.0	-38.09	-4.80	90.00	SportsCarDriver	0.00	61.91	E0_0
1	1	0.1	f_0.0	-38.08	-4.80	90.00	SportsCarDriver	0.03	61.92	E0_0
2	2	0.2	f_0.0	-38.07	-4.80	90.00	SportsCarDriver	0.10	61.93	E0_0
3	3	0.3	f_0.0	-38.05	-4.80	90.00	SportsCarDriver	0.21	61.95	E0_0
4	4	0.4	f_0.0	-38.01	-4.80	90.00	SportsCarDriver	0.40	61.99	E0_0
...
398475	398475	444.3	f_5.99	87.30	10.02	294.06	CautiousDriver	16.62	5.14	:J2_0_0
398476	398476	444.4	f_5.99	85.91	10.87	300.53	CautiousDriver	16.55	6.79	:J2_0_0
398477	398477	444.5	f_5.99	84.29	11.17	296.05	CautiousDriver	16.48	8.44	:J2_0_0
398478	398478	444.6	f_5.99	82.66	11.27	285.32	CautiousDriver	16.41	1.02	E8_0
398479	398479	444.7	f_5.99	81.03	11.25	274.53	CautiousDriver	16.34	2.65	E8_0

398480 rows × 16 columns



```
In [29]: # Merging the data using 'id' and 'type'
merged_data = pd.merge(fcd_df, emission_df, on=['time', 'id', 'type'])
```


```
In [30]: # Merging the data using 'id' and 'type'
merged_data = pd.merge(merged_data, features, on=['time', 'id', 'type'])
```

In [31]: merged_data


Out[31]:

	time	id	x_x	y_x	angle_x		type	speed_x	pos_x	l
0	0.0	f_0.0	-38.09	-4.80	90.00		SportsCarDriver	0.00	61.91	
1	0.0	f_1.0	-57.03	-37.67	33.37		CautiousDriver	0.00	75.68	
2	0.0	f_2.0	-33.33	-1.60	90.00		ElderlyDriver	0.00	66.67	
3	0.0	f_3.0	357.14	4.80	270.00		AggressiveLanechangingdriver	0.00	42.86	
4	0.0	f_4.0	241.15	21.32	215.86		AggressiveLanechangingdriver	0.00	98.24	
...	
398475	999.9	f_5.222	111.69	1.60	270.00		SportsCarDriver	0.00	102.38	
398476	999.9	f_5.223	245.35	4.80	270.00		InexperiencedDriver	12.57	55.00	
398477	999.9	f_5.224	158.14	4.80	270.00		SportsCarDriver	26.32	55.93	
398478	999.9	f_5.225	335.19	1.60	270.00		ElderlyDriver	8.05	64.81	
398479	999.9	f_5.226	367.06	4.80	270.00		ElderlyDriver	0.04	32.94	

398480 rows × 41 columns



In [32]: final = merged_data.drop(columns=['slope', 'eclass', 'route', 'waiting', 'lane_y',




In [33]: final

Out[33]:

	time	id	x_x	y_x	angle_x		type	speed_x	pos_x	a
0	0.0	f_0.0	-38.09	-4.80	90.00		SportsCarDriver	0.00	61.91	
1	0.0	f_1.0	-57.03	-37.67	33.37		CautiousDriver	0.00	75.68	
2	0.0	f_2.0	-33.33	-1.60	90.00		ElderlyDriver	0.00	66.67	
3	0.0	f_3.0	357.14	4.80	270.00		AggressiveLanechangingdriver	0.00	42.86	
4	0.0	f_4.0	241.15	21.32	215.86		AggressiveLanechangingdriver	0.00	98.24	
...
398475	999.9	f_5.222	111.69	1.60	270.00		SportsCarDriver	0.00	102.38	
398476	999.9	f_5.223	245.35	4.80	270.00		InexperiencedDriver	12.57	55.00	
398477	999.9	f_5.224	158.14	4.80	270.00		SportsCarDriver	26.32	55.93	
398478	999.9	f_5.225	335.19	1.60	270.00		ElderlyDriver	8.05	64.81	
398479	999.9	f_5.226	367.06	4.80	270.00		ElderlyDriver	0.04	32.94	

398480 rows × 17 columns




```
In [34]: from sklearn.preprocessing import LabelEncoder
# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Encode categorical columns
final['type'] = label_encoder.fit_transform(final['type'])
final['id'] = label_encoder.fit_transform(final['id'])
```

In [35]: `final.corr()`

Out[35]:

	time	id	x_x	y_x	angle_x	type	sp
time	1.000000	-0.027918	-0.011356	0.005042	0.014711	-0.014541	-0.0
id	-0.027918	1.000000	0.327327	0.409731	0.838318	0.065246	-0.2
x_x	-0.011356	0.327327	1.000000	0.386571	0.391870	0.017162	0.0
y_x	0.005042	0.409731	0.386571	1.000000	0.518102	0.037787	0.0
angle_x	0.014711	0.838318	0.391870	0.518102	1.000000	0.061652	-0.0
type	-0.014541	0.065246	0.017162	0.037787	0.061652	1.000000	0.0
speed_x	-0.047372	-0.203942	0.076041	0.001827	-0.078379	0.039832	1.0
pos_x	0.032194	0.066977	-0.114310	0.007374	-0.005626	0.013670	-0.1
acceleration_x	-0.021207	-0.047929	-0.026528	-0.005661	-0.045198	0.006287	-0.1
CO	-0.007374	0.047427	-0.039143	-0.006916	0.017719	0.085217	-0.3
fuel	-0.032181	-0.061821	-0.019635	0.002178	-0.014780	0.090033	0.2
noise	-0.060527	-0.173792	0.054022	0.003849	-0.077087	0.085638	0.6
angle	0.014711	0.838318	0.391870	0.518102	1.000000	0.061652	-0.0
lane_length	0.015000	0.036952	-0.226934	-0.112893	-0.090701	0.000577	-0.1
lane_change	-0.003851	-0.017214	0.012718	0.002876	0.000320	0.004922	0.0
vehicle_density	-0.004640	0.017209	0.014581	-0.016119	0.016353	-0.004523	-0.0
avg_speed_nearby_vehicles	-0.050949	-0.214982	0.067308	-0.013411	-0.087862	0.029756	0.9



Merging Mean, Variance and Standard Deviation data

```
In [36]: import pandas as pd

# Assuming 'df' contains the provided DataFrame

# Group by 'id' and 'time' and calculate the mean, std, and var for each column
mean_values = final.groupby(['id', 'type']).mean().reset_index()
std_values = final.groupby(['id', 'type']).std().reset_index()
var_values = final.groupby(['id', 'type']).var().reset_index()

# Merge the mean, std, and var DataFrames
result = pd.merge(mean_values, std_values, on=['id', 'type'], suffixes=('_mean', '_std'))
result = pd.merge(result, var_values, on=['id', 'type'], suffixes=('', '_var'))

# Display the resulting DataFrame with mean, std, and var for each column
print(result)
```

	id	type	time_mean	x_x_mean	y_x_mean	angle_x_mean	\
0	0	4	7.50	107.577550	-4.800000	90.000000	
1	1	2	16.00	87.693980	-3.451505	90.000000	
2	2	4	37.95	140.579252	-4.201869	90.000000	
3	3	1	364.15	74.287206	-4.800000	90.000000	
4	4	2	368.30	85.125269	-1.600000	90.000000	
...	
1266	1266	3	417.75	232.513088	7.357618	275.632353	
1267	1267	3	421.75	227.060992	5.506031	271.744313	
1268	1268	1	432.75	248.253301	10.328058	277.318204	
1269	1269	2	430.50	206.437138	10.499731	278.446229	
1270	1270	1	432.35	250.277097	5.909798	270.869234	

	speed_x_mean	pos_x_mean	acceleration_x_mean	CO_mean	...	\
0	23.237682	47.523974		1.463444	240.779139	...
1	12.496923	47.752943		0.657291	59.650535	...
2	20.085093	45.981869		0.996636	160.531729	...
3	12.905353	42.313765		0.627265	65.949529	...
4	12.671254	48.519319		0.670789	59.147742	...
...
1266	10.433412	43.509000		0.375382	60.537941	...
1267	10.203282	48.202863		0.555420	60.487786	...
1268	9.568932	44.238835		0.351092	58.010413	...
1269	11.411111	50.160101		0.534141	62.723064	...
1270	10.555524	50.547944		0.659032	68.910524	...

	pos_x	acceleration_x	CO	fuel	noise	\
0	949.568609	10.203543	41517.428541	23.190705	45.237318	
1	879.061975	0.211585	2212.185192	0.247460	9.977822	
2	838.181464	8.112813	29855.731260	15.921622	49.947717	
3	731.332891	0.185713	1779.774978	0.620957	14.549115	
4	886.990330	0.215334	2186.829440	0.219426	6.459293	
...	
1266	999.195548	1.002605	2031.932277	0.586912	8.121168	
1267	1201.809459	0.220230	2524.072358	0.308355	4.932657	
1268	824.095973	0.501527	2413.930779	0.651597	13.694340	
1269	999.236114	0.954994	1891.470833	0.476405	5.947707	
1270	963.310085	0.228923	2202.294516	0.403851	7.112195	

	angle	lane_length	lane_change	vehicle_density	\
0	0.000000	904.924682	0.056424	0.001447	
1	0.000000	773.581594	0.035555	0.001277	
2	0.000000	701.116068	0.040477	0.001276	
3	0.000000	639.056340	0.025846	0.001084	
4	0.000000	808.739493	0.031330	0.001287	
...	
1266	230.048384	962.355210	0.025846	0.000562	
1267	55.316171	1216.862500	0.033298	0.000967	
1268	310.465477	822.919064	0.019087	0.000535	
1269	346.527782	1123.479815	0.029484	0.000749	
1270	22.829946	1063.800649	0.027540	0.000791	

	avg_speed_nearby_vehicles
0	113.383921
1	39.974584
2	51.891929
3	53.400125

```

4          37.797974
...
1266      18.254119
1267      24.101123
1268      25.566286
1269      21.923327
1270      22.718884

```

[1271 rows x 47 columns]

```
In [37]: result.fillna(value=0, inplace=True) # Replaces NaN with 0
```

```
In [38]: result.columns
```

```
Out[38]: Index(['id', 'type', 'time_mean', 'x_x_mean', 'y_x_mean', 'angle_x_mean',
               'speed_x_mean', 'pos_x_mean', 'acceleration_x_mean', 'CO_mean',
               'fuel_mean', 'noise_mean', 'angle_mean', 'lane_length_mean',
               'lane_change_mean', 'vehicle_density_mean',
               'avg_speed_nearby_vehicles_mean', 'time_std', 'x_x_std', 'y_x_std',
               'angle_x_std', 'speed_x_std', 'pos_x_std', 'acceleration_x_std',
               'CO_std', 'fuel_std', 'noise_std', 'angle_std', 'lane_length_std',
               'lane_change_std', 'vehicle_density_std',
               'avg_speed_nearby_vehicles_std', 'time', 'x_x', 'y_x', 'angle_x',
               'speed_x', 'pos_x', 'acceleration_x', 'CO', 'fuel', 'noise', 'angle',
               'lane_length', 'lane_change', 'vehicle_density',
               'avg_speed_nearby_vehicles'],
              dtype='object')
```

```
In [52]: result['type'].value_counts()
```

```
Out[52]: type
3      406
2      265
4      249
1      239
0      112
Name: count, dtype: int64
```

Random Forest(ensemble Method)

```
In [39]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Assuming 'df' is your DataFrame containing the data

# Features (X) and target variable (y)
X = result.drop(['type', 'id'], axis=1) # Features
y = result['type'] # Target variable

# Split the data into training and testing sets (adjust test_size and random_s
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, rando

# Initialize the classifier (you can use any other classifier of your choice)
clf_rand = RandomForestClassifier(n_estimators=200, random_state=42)

# Fit the classifier on the training data
clf_rand.fit(X_train, y_train)

# Predict on the test data
y_pred = clf_rand.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
```

Accuracy: 0.7450980392156863

KNN-Model


```
In [40]: # Import necessary Libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

# Features (X) and target variable (y)
X = result.drop(['type', 'id'], axis=1) # Features
y = result['type'] # Target variable

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features by removing the mean and scaling to unit variance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Initialize the k-NN classifier
k = 5 # You can change this value as needed
knn = KNeighborsClassifier(n_neighbors=k)

# Fit the model on the training data
knn.fit(X_train, y_train)

# Predict the labels for the test set
y_pred = knn.predict(X_test)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy of the k-NN model: {accuracy:.2f}')
```

Accuracy of the k-NN model: 0.45

DNN Model

```

In [41]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras import regularizers
from tensorflow.keras.utils import to_categorical

# Assuming 'fcd_df' is your DataFrame containing the data

# Features (X) and target variable (y)
X = result.drop(['type', 'id'], axis=1) # Features
y = result['type'] # Target variable

# Convert y to categorical (one-hot encoded)
y = to_categorical(y)
num_classes = 5
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Define the Deep Neural Network model with L2 regularization and dropout
model = Sequential()
model.add(Dense(128, input_dim=X_train.shape[1], activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model.add(Dropout(0.5))
model.add(Dense(64, activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model.add(Dropout(0.5))
model.add(Dense(32, activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax')) # For multi-class classification

# Compile the model
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

# Fit the model on the training data
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))

# Evaluate the model on the test data
_, accuracy = model.evaluate(X_test, y_test)
print(f"Accuracy: {accuracy}")

```

```

Epoch 1/10
32/32 [=====] - 1s 7ms/step - loss: 3.7146 - accuracy: 0.2549 - val_loss: 3.2458 - val_accuracy: 0.4667
Epoch 2/10
32/32 [=====] - 0s 3ms/step - loss: 3.2216 - accuracy: 0.3947 - val_loss: 2.9941 - val_accuracy: 0.4902
Epoch 3/10
32/32 [=====] - 0s 3ms/step - loss: 2.9504 - accuracy: 0.4341 - val_loss: 2.7425 - val_accuracy: 0.4902
Epoch 4/10
32/32 [=====] - 0s 3ms/step - loss: 2.7408 - accuracy: 0.4675 - val_loss: 2.5296 - val_accuracy: 0.4902
Epoch 5/10
32/32 [=====] - 0s 3ms/step - loss: 2.5788 - accuracy: 0.4528 - val_loss: 2.3764 - val_accuracy: 0.4902
Epoch 6/10
32/32 [=====] - 0s 3ms/step - loss: 2.4209 - accuracy: 0.4587 - val_loss: 2.2328 - val_accuracy: 0.4902
Epoch 7/10
32/32 [=====] - 0s 3ms/step - loss: 2.2669 - accuracy: 0.4547 - val_loss: 2.1011 - val_accuracy: 0.4902
Epoch 8/10
32/32 [=====] - 0s 2ms/step - loss: 2.1251 - accuracy: 0.4715 - val_loss: 1.9885 - val_accuracy: 0.4902
Epoch 9/10
32/32 [=====] - 0s 3ms/step - loss: 2.0118 - accuracy: 0.4862 - val_loss: 1.8888 - val_accuracy: 0.4902
Epoch 10/10
32/32 [=====] - 0s 3ms/step - loss: 1.8991 - accuracy: 0.4902 - val_loss: 1.8035 - val_accuracy: 0.4902
8/8 [=====] - 0s 2ms/step - loss: 1.8035 - accuracy: 0.4902
Accuracy: 0.4901960790157318

```

Gaussian Naive Bayes

```

In [42]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import MinMaxScaler
import pandas as pd

# Assuming 'fcd_df' is your DataFrame containing the data

# Features (X) and target variable (y)
X = result.drop(['type', 'id'], axis=1) # Features
y = result['type'] # Target variable

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the MinMaxScaler and fit-transform the training data
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)

# Transform the test data using the same scaler
X_test = scaler.transform(X_test)

# Initialize the Gaussian Naive Bayes classifier
clf = GaussianNB()

# Fit the classifier on the training data
clf.fit(X_train, y_train)

# Predict on the test data
y_pred = clf.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

```

Accuracy: 0.6039215686274509

Best Model Performance on unseen data(New Simulation data)

```

In [19]: fcd_df_New = pd.read_excel("C:/Users/saiko/OneDrive/Desktop/699/Task-7/fcd-out
emission_df_New = pd.read_excel("C:/Users/saiko/OneDrive/Desktop/699/Task-7/em
features_New = pd.read_excel("C:/Users/saiko/OneDrive/Desktop/699/Task-7/feat


In [43]: # Merging the data using 'id' and 'type'
merged_data_New = pd.merge(fcd_df_New, emission_df_New, on=['time', 'id', 'type'])

```

```
In [44]: # Merging the data using 'id' and 'type'  
merged_data_New = pd.merge(merged_data_New, features_New, on=['time', 'id', 'type'])
```

Count Plot of Vehicle Types

```
In [55]: import seaborn as sns  
import matplotlib.pyplot as plt  
  
# Example of creating a count plot for the 'type' column  
sns.countplot(data=merged_data_New, x='type')  
plt.title('Count Plot of Vehicle Types')  
plt.xticks(rotation=45) # Rotate x-axis labels for better readability if needed  
plt.show()
```



```
In [45]: final_New = merged_data_New.drop(columns=['slope', 'eclass', 'route', 'waiting',
```



```
In [46]: from sklearn.preprocessing import LabelEncoder
# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Encode categorical columns
final_New['type'] = label_encoder.fit_transform(final_New['type'])
final_New['id'] = label_encoder.fit_transform(final_New['id'])
```

```
In [47]: import pandas as pd

# Assuming 'df' contains the provided DataFrame

# Group by 'id' and 'time' and calculate the mean, std, and var for each column
mean_values = final_New.groupby(['id', 'type']).mean().reset_index()
std_values = final_New.groupby(['id', 'type']).std().reset_index()
var_values = final_New.groupby(['id', 'type']).var().reset_index()

# Merge the mean, std, and var DataFrames
result_New = pd.merge(mean_values, std_values, on=['id', 'type'], suffixes=('_mean', '_std'))
result_New = pd.merge(result_New, var_values, on=['id', 'type'], suffixes=('_mean', '_std'))

# Display the resulting DataFrame with mean, std, and var for each column
print(result_New)
```

	id	type	time_mean	x_x_mean	y_x_mean	angle_x_mean	\
0	0	4	9.00	142.739171	-4.800000	90.000000	
1	1	3	17.95	52.987160	-3.209467	90.000000	
2	2	1	51.55	108.559641	-3.774586	90.000000	
3	3	2	365.00	80.617143	-4.800000	90.000000	
4	4	2	376.70	54.816736	-1.600000	90.000000	
...	
1267	1267	2	416.55	270.367073	5.593496	270.000000	
1268	1268	1	424.45	256.381689	4.590338	271.160473	
1269	1269	2	436.25	216.555180	11.507335	278.466437	
1270	1270	2	436.35	225.095988	6.357685	276.368056	
1271	1271	3	434.35	229.291532	6.304073	271.350685	

	speed_x_mean	pos_x_mean	acceleration_x_mean	CO_mean	...	\
0	20.700552	44.590884		1.043204	145.199448	...
1	12.304763	37.877456		0.543343	64.902781	...
2	12.019586	44.773785		0.387403	52.506436	...
3	12.947944	46.996620		0.697666	62.435331	...
4	11.289034	39.785326		0.503812	56.720209	...
...
1267	10.962724	46.412114		0.176992	63.994268	...
1268	10.374392	44.283986		0.362128	56.202095	...
1269	10.652605	48.867754		0.454311	56.270659	...
1270	10.213364	50.772130		0.485062	64.418981	...
1271	10.429556	48.939435		0.550726	67.486129	...

	pos_x	acceleration_x	CO	fuel	noise	\
0	1019.697413	6.982777	19547.554932	10.832321	35.822783	
1	846.849002	0.276305	1554.719180	0.635933	11.356579	
2	734.573604	0.495733	2030.685574	0.497010	6.973152	
3	819.476962	0.251244	2162.356078	0.387623	7.454172	
4	899.662123	0.486971	2198.392627	0.408777	16.535861	
...	
1267	809.955173	3.044184	2325.792560	0.646991	15.003130	
1268	946.908198	0.516791	2389.770224	0.634483	5.678417	
1269	909.139919	0.272772	2018.952948	0.477212	6.751585	
1270	912.561277	0.684037	2215.847815	0.527766	11.589522	
1271	1231.874856	0.451544	2336.980617	0.814688	6.293530	

	angle	lane_length	lane_change	vehicle_density	\
0	0.000000	884.953263	0.047514	0.001546	
1	0.000000	649.869537	0.028796	0.001054	
2	0.000000	651.669529	0.026936	0.001087	
3	0.000000	779.688419	0.030482	0.001084	
4	0.000000	620.317460	0.023007	0.001045	
...	
1267	0.000000	772.336051	0.027758	0.000405	
1268	27.584775	920.808450	0.026386	0.000619	
1269	354.483442	1077.393998	0.026299	0.000718	
1270	270.046015	1006.526309	0.027090	0.000551	
1271	31.143291	1234.263667	0.027540	0.001025	

	avg_speed_nearby_vehicles
0	65.350769
1	40.054214
2	44.091203
3	35.755170


```
4          41.627853
...          ...
1267       19.069430
1268       27.404635
1269       21.426673
1270       29.590376
1271       20.784419
```

[1272 rows x 47 columns]

```
In [48]: result_New.fillna(value=0, inplace=True) # Replaces NaN with 0
```

```
In [53]: result_New['type'].value_counts()
```

```
Out[53]: type
3      384
4      260
1      244
2      239
0      145
Name: count, dtype: int64
```

```
In [49]: import pandas as pd

# Assuming df is your DataFrame
# Shuffle the rows using sample() function
shuffled_df = result_New.sample(frac=1, random_state=42) # frac=1 shuffles the entire DataFrame

# Reset the index if needed
shuffled_df.reset_index(drop=True, inplace=True)
```

Random Forest Evaluation

```
In [51]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Assuming 'df' is your DataFrame containing the data

# Features (X) and target variable (y)
X = shuffled_df.drop(['type', 'id'], axis=1) # Features
y = shuffled_df['type'] # Target variable

# Predict on the test data
y_pred = clf_rand.predict(X)

# Calculate accuracy
accuracy = accuracy_score(y, y_pred)
print(f"Accuracy: {accuracy}")
```

Accuracy: 0.7130503144654088

In []: